ITF22519: Introduction to Operating Systems

Fall Semester, 2022

Lab5: Processes in Linux

Submission Deadline: September 27$^{\text{th}}$, 2022 23:59

You need to get at least 50pts to pass this lab assignment.

In this lab, you will learn how to create a new process in Linux Operating System and do some practices with several system calls such as `fork`, `sleep`, `exec`, `wait`, and `kill`. Before you start, pull the new lab:

```
$ cd OS2022/labs
$ git pull main main
$ cd lab5
```

1 About Unix Processes

- When a system is booted, the first user space process is `systemd` or `/sbin/init` depending on your Linux distribution. This process has proccess id (PID) of 1. It will launch startup scripts and eventually login prompts. If you do a `ps -el`, you should see that process 1 is `systemd`. It is the ancestor of all other user processes on the system.

- When you login or start a terminal, a process for the shell is started. The shell will then launch other processes, which will be children of the shell. If the parent process dies, `systemd` will adopt the orphaned processes.

- The *state* of a Unix process is shown as the second column of the process table (viewed by executing the `ps` command). Some of the states are R: running, W: waiting, S: sleeping, Z: zombie.

Figure 1: Example of Unix processes

Whenever a program or a command executes, there is a process. A process can be run in foreground or background.

- Foreground: Every process when started runs in foreground by default. It receives input from the keyboard and then sends output to the screen. When a command/process is running in the foreground, it takes a lot of time and no other processes can be run because the prompt would not be available until the program finishes processing and comes out.

- Background: A process runs in the background without keyboard input and waits till keyboard input is required. Thus, other processes can be done in parallel with the process running in the background since they do not have to wait for the previous process to be completed. Adding suffix with an ampersand & starts the process as a background process.

1.1 Process Table

The program *print_pid.c* prints out its process ID (PID) and its parent process ID (PPID), and then sleeps for 30 seconds.

- Make the executable file output named *print_pid* for *print_pid.c*

- Run *print_pid* twice, both times as a background process.

- In the meantime, press `Ctrl+C` to terminate the process - see the commands below.

- When you see the message "I am awake", the proccess is finished and no longer show up on the screen.

```
$ ./print_pid &
Ctrl+C
$ ./print_pid &
Ctrl+C
```

Now:

- Do `ps -el` to view all processes in the sysem

- Do `ps -l` to view only processes running on your terminal

- Do `ps -l | less` to pipe output to the less

- Do `ps -l > output` to direct output the the file named `output`

The first line from a the command `ps -l` is:

```
F S    UID    PID PPID  C PRI   NI ADDR    SZ WCHAN    TTY    TIME CMD
```

A short description of those fields can be found in this link.

1.2 Killing a process

You have known how to find the PID of a process. If you want to terminate an unwanted process, you can use the command:

```
$ kill PID
```

If the process refuses to be killed, type:

```
$ kill -9 PID
```

Here `PID` is not `PID` but the ID of the process you want to kill and you have permission to kill. You might also try `killall` to kill all processes. Do `man killall` to know more about `killall`.

2 System call

In a computer, the **user space** is what most users interact with when using the computer. It is a portion of memory with restricted permission and access. Under the user space is the **kernel space** - the actual operating system (OS). Kernel space has complete access to everything on the system.

Because of privileged rights, it is not wise to give any user the access to the kernel space. Therefore, there is a restricted and well-defined interface the between user space and the kernel space: the **system calls**. When a system call is called, the program execution stops and the execution is switched to the OS address space. The parameters are passed to the OS through a pre-determined register passing scheme. The OS will then check if the program has the permissions to perform the requested operations. If yes, the task is completed.

2.1 The fork() system call: Creating a new process

The `fork()` system call is used to create a new process which becomes the child process of the calling process. If the `fork()` does not return a negative value, the creation of the child process is successful. If `fork()` returns zero, the process is the child process. Otherwise, i.e, `fork()` returns a positive, the process is the parent process.

```
#include <sys/types.h>
#include <unistd.h>
// ...
pid_t child;
child = fork();
// ...
```

In this code snippet, the `fork()` creates a new variable called `child`. The `child` is of type `pid_t` defined in `sys/types.h`. In addition, the `fork()` is defined in the header `unistd.h`. Therefore, you have to add `sys/types.h` and `unistd.h` on the top of your code.

The child process is the copy of the parent process. After spawning a child process, all of the statements after `fork()` system call are executed by both the parent and child processes. However, both processes are now separate and changes in one do not affect the other. The code fragment below calls fork() once, which results in the creation of one child process. Then, each process prints out its id and the parent's ID.

```
int main() {
    fork();
    printf("Process %d's parent process ID is %d\n", getpid(), getppid());
    return 0;
}
```

The output of this program can be:

```
Process 14427's parent process ID is 6891
Process 14428's parent process ID is 14427
```

Here, the ID of the main function is 14427, the child ID is 14428. The the ID of bash shell is 6891. The process tree of this program is showed below.

Figure 2: A process tree

The program *fork_ex1.c* spawns new processes with two `fork()` in the `main()`. The output of the program can be:

```
Process 35875's parent process ID is 33183
Process 35876's parent process ID is 35875
Process 35877's parent process ID is 35875
Process 35878's parent process ID is 35876
```

Figure 3: A possible output of *fork_ex1.c*

Figure 4: Corresponding process tree

Task 1: The following program prints different messages in the child and parent processes. Complete the condition of the if statements in this program and save the code in a *fork_ex2.c* file.

```c
#include <sys/types.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

int main() {
    pid_t pid;
    pid = fork();
    if ( ) { // ADD YOUR CODE HERE
```

5

```
        printf("This is the child process");
        printf("The child process ID is %d\n", getpid());
        printf("The child's parent process ID is %d\n", getppid());
    }
  else if ( ){  // ADD YOUR CODE HERE

        printf("This is the parent process");
        printf("The parent process ID is %d\n", getpid());
        printf("The parent's parent process ID is %d\n", getppid());
 }
 else {
        perror("fork"); // fork() fails; handle error here
        exit(-1);
 }
 sleep(2);
 return 0;
}
```

2.2   The wait() system call: Waiting on a child process

The child process and its parent process are at their own places. Therefore, they execute independently and one can finish before the other. In some situations, this might not be the desired behavior. There are two system calls which guarantee that the parent process waits for its child process to complete. These system calls are `wait()` and `waitpid()`.

`wait():`   This function blocks parent process until the child process (whose PID stored in `child`) exits or a signal is received. After child process terminates, parent process continues its execution after `wait()` system call instruction. The snippet of code for `wait()` is as follows:

```
#include <sys/types.h>
#include <sys/wait.h>
...

int status = 0;
....

pid_t pid;
pid = fork();
if ( pid == 0){
// This is the child process.
// Do something here
} else if ( pid >0 ) {
// This is the parent process.
// Wait for the child to finish
        wait(&status);
// Do something else here
        printf("The child is done, status is: %d\n", status);
        return 0;
} else {
        perror("fork");
        exit(-1);
```

```
}
```

Task 2: Complete the above code, name the program as *fork_ex2.c* and commit the file *fork_ex2.c* into your Github.

waitpid(): If the parent process has multiple children, the knowedge of a specific child process' termination is of importance. In this case, waitpid() system call should be used instead of wait(). The waitpid() specifies which child process the parent process is waiting for. The snippet of code for using waitpid() is as follows:

```c
#include <sys/types.h>
#include <sys/wait.h>
...
int status = 0;
...
child = fork();
if(child == 0){
// This is the child process.
// Do something here
}
else if(child > 0){
// This is the parent process.
// Wait for a specific child to finish
    waitpid(child, &status, 0);
    printf("child process is done, status is: %d\n", status);
// do something else here
    return 0;
}
else {
    perror("fork");
    exit(-1);
}
```

For more information about wait() and waitpid(), use man page or google it.

Task 3: Compile and execute the programs *no_wait.c*, *wait.c*, *waitpid.c*.

- Redirect the output of each program to a file.

- Concatenate above three output files and include the result into your report.

- Explain the difference of the three programs.

2.3   The execl() system call: Making a process run another program

The fork() system call is useful for splitting a process into two, but it is not very useful to execute new programs. To start a new program, a different system call which is execve() family of system call must be used.

There are a number of available functions defined in execve() family. In this lab, we look at execl function. In our program, we will use execl in the main so that our main will run another defined program. The syntax for execl is as follows:

```
int execl(const char *pathname, const char *arg, ..., NULL);
```

The `execl` function is defined in the header `unistd.h`, so remember to add the header in your code. Look at the following code snippet:

```
#include <stdio.h>
#include <unistd.h>

int main() {
    execl("/bin/ls", "ls", NULL);
}
```

In this example, `/bin/ls` is the path of the executable bin file. `ls` is the program you want to execute. If there is any errors, the `execl` returns `-1`. Otherwise, it returns nothing.

2.4   The kill() system call: Sending a signal to another process

It is not that the `kill()` system call is used to terminate another process. It is used to send all kinds of signal to a process. Snippet for using `kill()` system call is as the following:

```
#include <sys/types.h>
#include <signal.h>
int kill(pid_t pid, int sig);
```

3   Exercises

3.1   Exercise 1 (30 points)

In the subsection 1.1 with the process `print_pid`:

- Take the screenshoot for the output of the command `ps -l`. Then, point out the following fields: `process name`, `Process state`, `process ID`, `ID of the parent process`, `time being created`.

- What is the process that spawned your *print_pid.c* processes.

- Kill the two *print_pid.c* processes and include the screenshoot of the command in your lab5 report.

- Do `ps -el` to see how the process table changes before and after your killing processes. Include only relevant lines for your programs in your report.

3.2   Exercise 2 (40 points)

Compile and execute the program *fork_ex3.c*.

- Include the screenshoot of the output in your lab5 report.

- Draw the process tree and label each process with its PID.

- If there are **n** fork() system call in the code, how many processes are created? Why?

- Remove `sleep` statement in the code. Run the code and explain what happens.

Note: **n** is an integer number.

3.3 Exercise 3 (30 points)

Refer to section 2.3 about `execl` function.

- Write a C program with which when running its executable bin file, the program behaves as if you type one of the following commands in your terminal: `ls, pwd, cat, ps, and, -l`.

- Include the screenshoot of your run and the output for each case in your report.

- Use `printf` function to print out a message after `execl` function. On which condition, the message is printed?

Note: You may want to use the following paths: `bin/ls`, `bin/cat`, `bin/ps`, `bin/pwd`.

4 What To Submit

Complete the exercises in this lab. Then, put all of files into the **lab5** directory of your repository. Make a report for each exercise. After that, run `git add .` and `git status` to ensure the file has been added and commit the changes by running `git commit -m "Commit Message"`. Finally, submit your files to GitHub by running `git push`. Check the GitHub website to make sure all files have been submitted.